

AI Fact-Checking System

Group ID
COMP90042 Project 2026

Abstract

We present a resource-aware climate fact-checking system optimized for limited compute, with a standard Colab-style runtime as the concrete constraint. The pipeline stages evidence processing: sparse fusion preserves broad recall, embedding plus factual-cue reranking compresses the pool, a bounded cross-encoder pass selects submitted evidence, and a TF-IDF logistic classifier predicts the verdict from structured evidence text. This Colab-feasible design obtains 0.2021 evidence F-score, 0.5000 claim accuracy, 0.2879 harmonic mean, and 0.4659 classifier macro-F1.

1 Introduction

Climate claims are short, compressed, and evidence-sensitive; a useful fact-checker must retrieve evidence as well as predict a label. Practical systems also face limited compute, so full-corpus neural scoring is often unavailable. We instantiate this constraint with a Colab-style runtime: no external services, saved checkpoints, or precomputed ranked artifacts. The problem is therefore a three-way trade-off among evidence recall, verdict accuracy, and runtime.

Draft note: I collected and grouped the following papers for teammates to continue the related-work writing.

Existing work can be organized into four areas:

1. **Evidence-Based Fact Checking** FEVER frames fact checking as joint evidence retrieval and verdict prediction (Thorne et al., 2018); CLIMATE-FEVER adds climate-specific domain and retrieval challenges (Diggelmann et al., 2020).
2. **Large-Scale Evidence Retrieval** BM25 is fast and scalable (Robertson et al., 1994); reciprocal rank fusion combines ranked lists without supervision (Cormack et al., 2009).

These methods suit high-recall candidate generation but remain lexical.

3. **Neural Semantic Reranking** Dense embeddings help when lexical overlap is weak (Karpukhin et al., 2020; Reimers and Gurevych, 2019); cross-encoders model claim–evidence interaction directly (Nogueira and Cho, 2019), but their cost requires sparse prefiltering.
4. **Evidence Aggregation and Factual Cue Fusion** Evidence aggregation studies passage combination before classification (Zhou et al., 2019). Hybrid systems show lexical, semantic, and shallow factual signals can be complementary (Lee et al., 2018; Padia et al., 2018).

Our design is literature-aligned but resource-driven. It first builds a broad unsupervised sparse pool from word, character, structured, and PRF gates. It then compresses this pool with MiniLM similarity plus factual features, applies a bounded cross-encoder only after prefiltering, and classifies from rank-aware, cue-aware evidence text. Contributions: (i) a staged pipeline for the performance–runtime trade-off, (ii) multi-gate sparse retrieval, (iii) multi-scale reranking that fuses semantic and shallow factual signals, and (iv) diagnostics over sparse gates, fusion, evidence selection, and classifier sensitivity.

2 Method

2.1 Pipeline Overview

Let $\mathcal{C}$ be claims and $\mathcal{E}$ evidence. The output is a verdict and a small evidence set per claim. Since scoring all $(c, e) \in \mathcal{C} \times \mathcal{E}$ pairs is infeasible, we use a staged pipeline:

$$\mathcal{E} \rightarrow \mathcal{A}(c) \rightarrow \mathcal{B}(c) \rightarrow \mathcal{D}(c),$$

where $\mathcal{A}(c)$ is a broad sparse pool, $\mathcal{B}(c)$ a reranked context pool, and $\mathcal{D}(c)$ the submitted evidence.




Figure 1: Staged pruning keeps recall while reducing pair count. The candidate and context pools retain most recoverable evidence without full-corpus neural scoring, which would require about 186 million development pairs.




Figure 2: Embedding plus factual cues is the best top-64 selector, reaching 57.8% macro recall above sparse cutoff and cross-encoder-plus-cue variants.

Cheap stages maximize recall; expensive semantic models operate only after pruning. The four components are multi-gate sparse retrieval, multi-scale reranking, bounded cross-encoder selection, and structured-evidence classification.

2.2 Multi-Gate Sparse Candidate Retrieval

The sparse stage maximizes recall. A gate is a feature construction, not a new scoring formula: the four gates differ in analyzer $a_g(\cdot)$ and vocabulary V_g , while sharing BM25. Let $q_g(c)$ and $x_g(d)$ be the resulting count vectors.

Word gate. Word terms give the direct lexical view: content-word overlap raises scores and IDF discounts frequent, weakly discriminative words.

Character gate. Character n-grams capture partial overlap, inflection, hyphenation, and spelling variation missed by exact words.

Structured gate. Structured cues keep numbers, years, salient capitalized terms, and negation markers. Standalone recall is low, but the cues add factual constraints that lexical views may lose.

Pseudo-relevance feedback gate. PRF expands the word query using high-ranked feedback evidence. If $\mathcal{F}(c)$ is the feedback set,

$$q_{\text{prf}}(c) = q_{\text{word}}(c) + \eta \sum_{d \in \mathcal{F}(c)} w(c, d) x_{\text{word}}(d),$$

where $w(c, d)$ downweights lower-ranked feedback documents. PRF adds related terms absent from the original claim while staying in the sparse word space.

Shared BM25 interaction. After feature construction, each gate applies the same BM25 interaction. For term t in document d ,

$$\text{idf}_t = \log \left(\frac{N - df_t + 0.5}{df_t + 0.5} + 1 \right),$$

$$\text{BM25}_{d,t} = \text{idf}_t \frac{(k_1 + 1) f_{d,t}}{f_{d,t} + k_1 (1 - b + b |d| / |d|)}.$$

IDF is part of BM25, not a separate gate; it discounts high-frequency features. The query side remains a count vector, giving the sparse dot product

$$s_g(c, d) = q_g(c)^\top \text{BM25}_g(d).$$

Thus g chooses the representation, and BM25 weights it. The notebook uses SciPy/NumPy sparse matrices, so document-frequency counts, BM25 reweighting, and query–document dot products run in optimized C/C++ kernels and vectorized array operations.

Rank fusion. Each gate produces a ranked list. The fifth sparse view is their weighted RRF fusion:

$$S_{\text{RRF}}(c, d) = \sum_{g \in G} \frac{\lambda_g}{K_{\text{RRF}} + r_g(c, d)},$$

where missing documents contribute zero. We treat λ_g as a reliability prior: stronger standalone gates receive larger weights, while weak but complementary signals can retain small non-zero weights. Structured recall is low alone, yet ablation shows it improves fused top-500 recall; RRF also avoids cross-gate score calibration.




Figure 3: Sparse fusion gives the best top-500 recall and retains structured cues for complementarity. Character matching is the strongest single source; structured cues are weak alone, but removing them lowers fused recall under a controlled ablation, so the final system keeps a small structured weight.

Figures 1 and 3 justify the candidate stage: fusion improves top-500 recall over any single source, and staging avoids full-corpus neural scoring.

2.3 Multi-Scale Feature Fusion

Reranking fuses dense similarity, sparse source rank, word overlap, numeric agreement, negation compatibility, and length shape. MiniLM embeddings are mean-pooled and ℓ_2 -normalized, so cosine similarity is an inner product:

$$s_{\text{emb}}(c, d) = h(c)^\top h(d).$$

The shallow factual features are summarized as

$$\phi(c, d) = [\phi_{\text{rank}}, \phi_{\text{lex}}, \phi_{\text{num}}, \phi_{\text{neg}}, \phi_{\text{len}}, s_{\text{emb}}].$$

where the components encode source rank, lexical overlap, number match, negation compatibility, length mismatch, and embedding similarity. A lightweight gradient-boosted matcher estimates $p_{\text{feat}}(c, d)$, and the context score is

$$S_{\text{ctx}}(c, d) = \alpha \tilde{s}_{\text{emb}}(c, d) + (1 - \alpha) p_{\text{feat}}(c, d),$$

where $\tilde{s}_{\text{emb}}$ is claim-wise normalized. Embeddings handle semantic relatedness; factual features guard against number, negation, and weak lexical-grounding errors.

2.4 Bounded Neural Evidence Selection

The submitted evidence stage uses a stricter operating point than the classifier context. Because cross-encoders jointly encode claim and evidence but are expensive, we apply them only after sparse/embedding prefiltering and normalize their scores within each claim. The final evidence score




Figure 4: Top-64 fusion raises the retained-candidate score distribution and reduces dispersion. In the Colab diagnostic, fused scores have mean/std = 0.682/0.106, compared with 0.525/0.109 for embeddings and 0.532/0.181 for factual cues.

fuses cross-encoder confidence, embedding confidence, embedding rank, and sparse source rank:

$$S_{\text{sub}}(c, d) = \beta_1 \tilde{s}_{\text{CE}} + \beta_2 \tilde{s}_{\text{emb}} + \frac{\beta_3}{\tau + r_{\text{emb}}} + \frac{\beta_4}{\tau + r_{\text{src}}}.$$

Coefficients are fixed. Candidates outside the cross-encoder subset receive no CE contribution but can still be selected by embedding and source-rank terms, bounding neural cost without discarding useful sparse candidates.

Figures 2, 4, and 5 justify multi-scale fusion: factual cues are not standalone rules, but complements that improve semantic ranking and stabilize the retained candidate pool.




Figure 5: Factual cues preserve or improve semantic ranking. Gains are small for the strong cross-encoder but clearer for embedding inner product, supporting feature fusion over a purely semantic selector.

2.5 Structured Evidence Classification

The classifier predicts from a bounded ranked evidence context. Rather than an unordered passage block, it constructs

$$x_c = [\text{claim}; z_1, z_1, \dots, z_m, z_m, z_{m+1}, \dots, z_L],$$

where top evidence is repeated to raise TF-IDF weight. Each evidence line also gets a rank prefix plus word-overlap, numeric-match, and negation-match cue tokens. These are learnable tokens, not rules. Let v_c be the TF-IDF vector; a one-vs-rest logistic model learns one binary classifier per label:

$$P(y = k | c) = \sigma(\theta_k^\top v_c + b_k).$$

The largest decision score gives the label. This keeps classification cheap while preserving retrieval rank and factual signals.

Stage	F-score	Macro recall	Hit-any
Sparse top-500	0.0083	0.6754	0.9026
Context top-64	0.0517	0.5784	0.8636
Submitted top-3	0.2021	0.2355	0.4675

Table 1: Retrieval metrics from the self-generated notebook artifacts.




Figure 6: Recall-vs-K curves show the pruning trade-off behind the selected operating points. The sparse pool keeps broad top-500 recall, the reranked pool preserves most recoverable evidence at top-64, and the submitted top-3 is the strict evidence set used for scoring.

3 Results

The final notebook was run in Colab with GPU enabled. Course metrics are evidence F-score F , label accuracy A , and H-score $H = 2FA/(F+A)$. Tables 1–2 report retrieval operating points and reproduced development results.

Figures 6 and 7 show the final retrieval-classification behavior. Remaining classification errors mainly confuse SUPPORTS, REFUTES, and NEI; DISPUTED has only 18 development examples.




Figure 7: Classifier diagnostics. Raw counts show total error mass, row-normalized confusion controls for class support, and per-class precision/recall/F1 summarizes the classification report.

Metric	Value
Evidence Retrieval F-score (F)	0.2021
Claim Classification Accuracy (A)	0.5000
Harmonic Mean of F and A (H)	0.2879
Top-3 Macro Recall	0.2355
Classifier Macro-F1	0.4659

Table 2: Final Colab-reproduced development result. F , A , and $H = 2FA/(F + A)$ are official course metrics; the last two rows summarize retrieval and classification behavior.

Table 3 justifies the classifier input: enhanced context improves over plain context at $C = 0.25$ (0.4935/0.4544 vs. 0.4351/0.3762), and the diagnostic best is $C = 0.125$ (0.5195/0.4841). The reported Colab run is Table 2: $F = 0.2021$, $A = 0.5000$, $H = 0.2879$.

Context format	C	Acc.	Macro-F1
Plain	0.25	0.4351	0.3762
Rank-aware	0.25	0.4351	0.3749
Top-weighted	0.25	0.4416	0.3938
Rank-weighted	0.25	0.4351	0.3882
Enhanced	0.25	0.4935	0.4544
Enhanced	0.125	0.5195	0.4841

C	Acc.	Macro-F1
0.125	0.5195	0.4841
0.25	0.4935	0.4544
0.5	0.4740	0.4263
1.0	0.4675	0.4184
2.0	0.4416	0.3825
4.0	0.4545	0.3913
8.0	0.4481	0.3925

Table 3: Classifier ablations. Rank-aware, top-weighted, factual-cue-enhanced context improves over plain concatenation; the enhanced classifier peaks at $C = 0.125$ in the diagnostic sweep.

References

- Gordon V. Cormack, Charles L. A. Clarke, and Stefan Buettcher. 2009. Reciprocal rank fusion outperforms Condorcet and individual rank learning methods. In *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 758–759.
- Thomas Diggelmann, Jordan Boyd-Graber, Jannis Bulian, Massimiliano Ciaramita, and Markus Leippold. 2020. CLIMATE-FEVER: A dataset for verification of real-world climate claims. *arXiv preprint arXiv:2012.00614*.
- Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 6769–6781.
- Nayeon Lee, Chien-Sheng Wu, and Pascale Fung. 2018. Improving large-scale fact-checking using decomposable attention models and lexical tagging. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 1133–1138.
- Rodrigo Nogueira and Kyunghyun Cho. 2019. Passage re-ranking with BERT. *arXiv preprint arXiv:1901.04085*.
- Ankur Padia, Francis Ferraro, and Tim Finin. 2018. Team UMBC-FEVER: Claim verification using semantic lexical resources. In *Proceedings of the First Workshop on Fact Extraction and VERification*, pages 161–165.
- Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, pages 3982–3992.
- Stephen E. Robertson, Steve Walker, Susan Jones, Micheline Hancock-Beaulieu, and Mike Gatford. 1994. Okapi at TREC-3. In *Proceedings of the Third Text REtrieval Conference*, pages 109–126.
- James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. 2018. FEVER: A large-scale dataset for fact extraction and verification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 809–819.

Jie Zhou, Xu Han, Cheng Yang, Zhiyuan Liu, Lifeng Wang, Changcheng Li, and Maosong Sun. 2019. GEAR: Graph-based evidence aggregating and reasoning for fact verification. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 892–901.